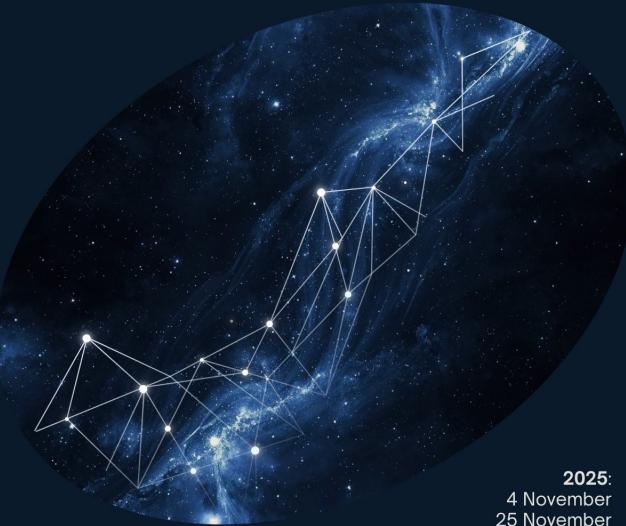


Introductory Deep Learning Lectures



2025:

4 November
25 November

2026:

20 January
10 February
17 March
21 April
19 May
16 June

Topics:
Neural Networks & Applications
Bayesian Deep Learning
Foundation Models
Self-supervised Learning

Presented by the
IACDEEP
Research Group

1h lectures
11 AM, IAC Aula

Instituto de Astrofísica de Canarias
C/ Vía Láctea, s/n 38205 La Laguna
Contact: iacdeeplectures@gmail.com



Intro to Neural Networks

Introductory Deep Learning lectures (IACDEEP)

Carlos Westendorp & Marc Huertas

https://github.com/cwestend/IACDEEP_introNN

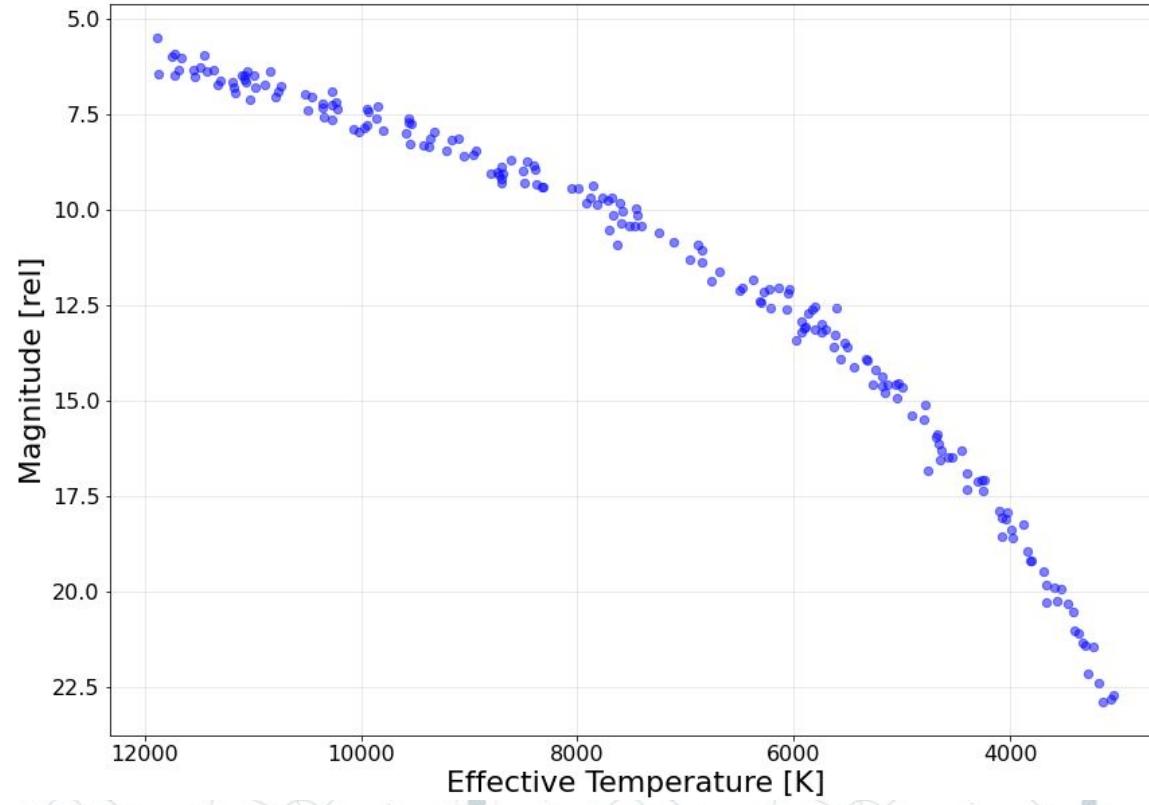


Course schedule

1. **Intro to NN (today)**
2. **Computer Vision using CNNs (25th November)**
3. Statistical deep learning
 - a. Bayesian statistics
 - b. Neural density estimators
 - c. Simulation-based inference
4. NNs for sequences / time series
5. NNs for unstructured data: Graph NNs
6. Foundational Models / self supervised learning

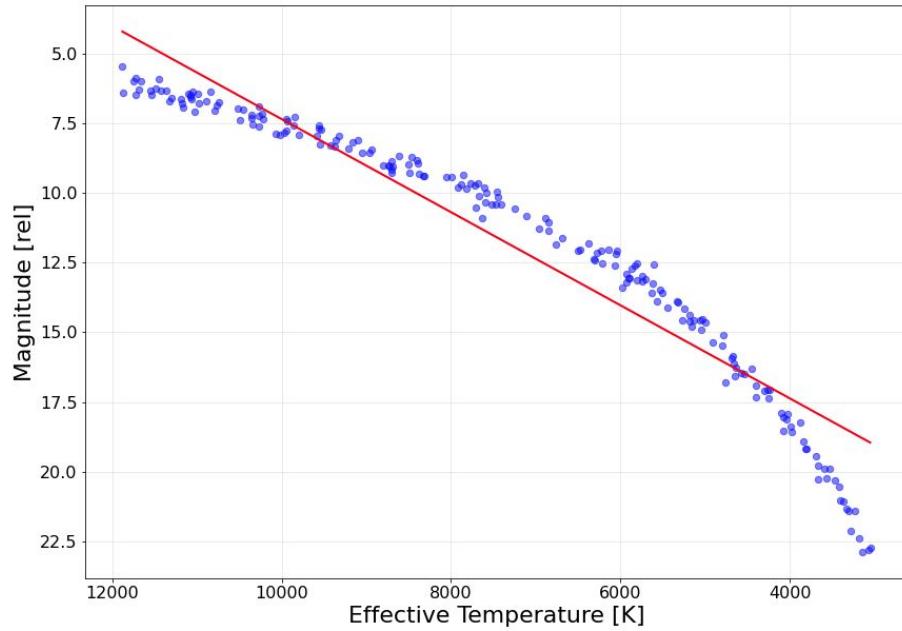


General problem




General problem

From physical insight



Data driven



Data driven

- ◎ No suitable **physical model** available: **accuracy**
- ◎ Physical Model **too complex** or **dataset too large**, minimisation difficulty: **speed**
- ◎ Possible **hidden information** in the data (beyond usual summary statistics): **discovery**



Astrophysics: **large** and **complex** datasets

Supervised learning

Given a dataset with **known labels** - find a function that can assign (predict) labels for an unlabelled dataset using a set of features (measurements)

Training Set

$$(\vec{x}_1, \vec{x}_2, \vec{x}_3, \dots, \vec{x}_n)$$



Features: colors, fluxes, spectral indexes (**Teff**)

$$(\vec{y}_1, \vec{y}_2, \vec{y}_3, \dots, \vec{y}_n)$$



Labels: morphology, object type, redshift (**magnitudes**)



Supervised learning

Given a dataset with **known labels** - find a function that can assign (predict) labels for an unlabelled dataset using a set of features (measurements)

Training Set

$$(\vec{x}_1, \vec{x}_2, \vec{x}_3, \dots, \vec{x}_n)$$

$$(\vec{y}_1, \vec{y}_2, \vec{y}_3, \dots, \vec{y}_n)$$



$$f_W(\vec{x}) = \vec{y}$$



Supervised learning

Unlabelled Set

$$(\vec{x}_1', \vec{x}_2', \vec{x}_3', \dots, \vec{x}_n')$$



$$(\vec{y}_1', \vec{y}_2', \vec{y}_3', \dots, \vec{y}_n')$$

Training Set

$$(\vec{x}_1, \vec{x}_2, \vec{x}_3, \dots, \vec{x}_n)$$

$$(\vec{y}_1, \vec{y}_2, \vec{y}_3, \dots, \vec{y}_n)$$



$$f_W(\vec{x}) = \vec{y}$$

?



Supervised learning

General Goal: Find a (**non-linear**) function that outputs the correct class / value (y) for a given input object:

$$f_W(\vec{x}) = \vec{y}$$

↓ ↓
Parameters (can be large!) Features

Minimization problem: find W such prediction error is minimal over all unseen vectors



Minimize the loss

1. Define a **Loss function**

$$loss(f_W(), \vec{x}_i, \vec{y}_i)$$

- for example: **MSE** loss

$$(f_W(\vec{x}_i) - \vec{y}_i)^2$$

2. **Minimize the empirical risk** with optimization

$$R_{\text{empirical}}(W) = \frac{1}{N} \sum_i^N loss(f_W(), \vec{x}_i, \vec{y}_i)$$



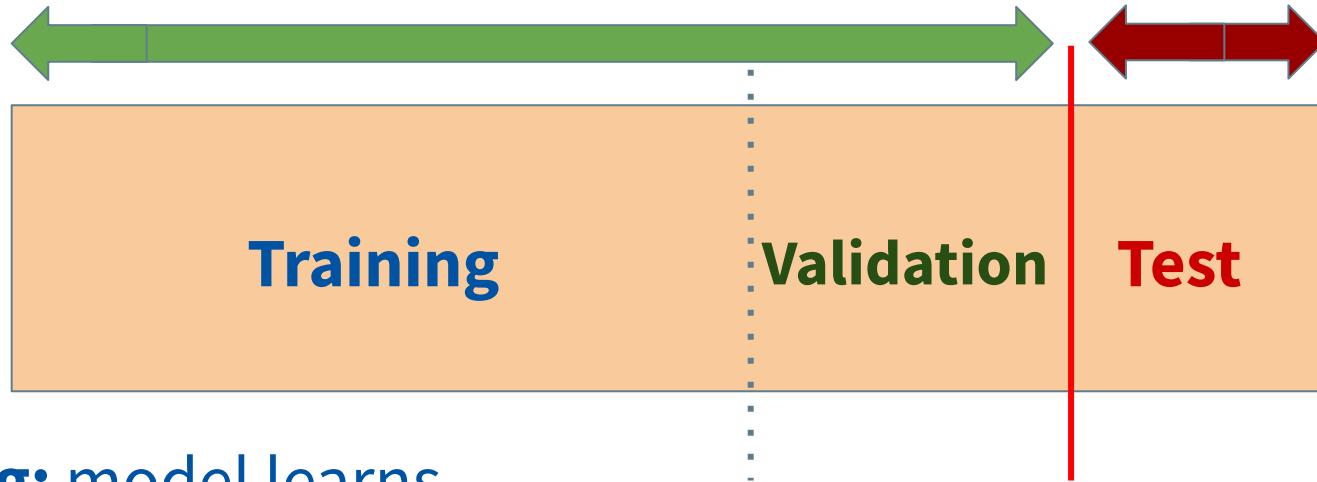
Minimize the loss $R_{empirical}(W) = \frac{1}{N} \sum_i^N loss(f_W(), \vec{x}_i, \vec{y}_i)$

ALL “GALAXIES IN THE UNIVERSE”

OBSERVED DATASET



In practice: **split data** (need enough!)



Training: model learns

Validation: monitor learning (overfitting)

Test: validity check (not used in training!)

Minimisation problem

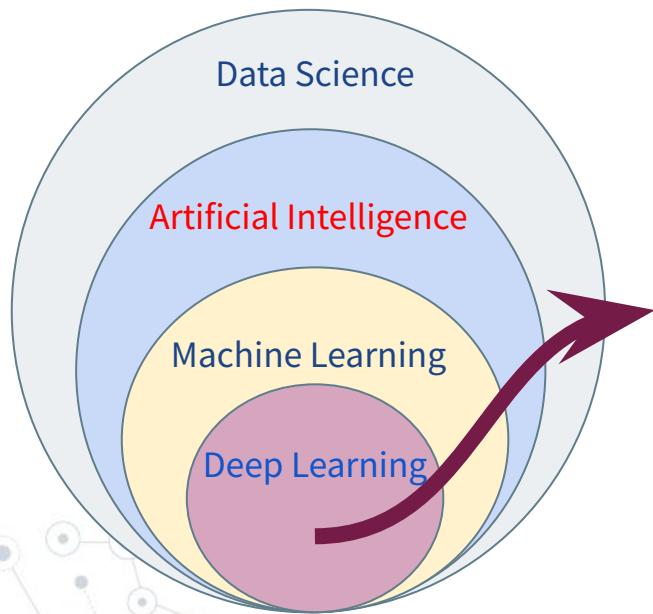
We need:

1. A **Loss function** (something to minimize)
2. Minimization (optimization) **algorithm**

common to **all** Machine Learning algorithms

Machine Learning and Deep Learning

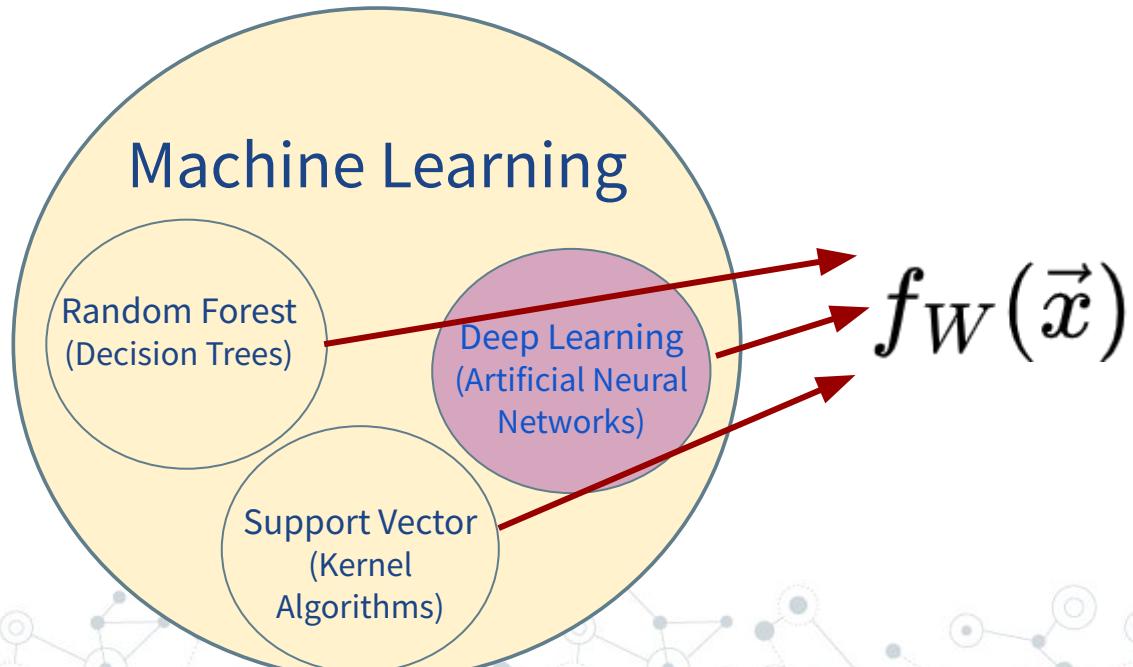
Deep Learning *within* Machine Learning



- Not only LLMs (ChatGPT, Gemini...)
- Healthcare (cancer detection)
- Autonomous vehicles
- Climate (forecasting, monitor)
- Speech & Audio (translating)
- Finance (fraud detection)
- ...


Machine Learning and Deep Learning

Difference is the function used (sets optimization/loss)



Machine Learning and Deep Learning

Deep Learning uses Artificial Neural Networks





Machine Learning

Random Forest
(Decision Trees)

Deep Learning
(Artificial Neural Networks)

Support Vector
(Kernel
Algorithms)





¿Why Deep Learning?



Deep Learning uses **Artificial Neural Networks**

Neural Network origins

- ◎ Automata = “that operates by itself” ancient concept (China, Islam, Greece - 300 BC)
- ◎ 1950: **Alan Turing** published “Computing Machinery and Intelligence” - Turing Test (called Imitation Game)
- ◎ 1956: **John McCarthy** workshop in Dartmouth about “**Artificial Intelligence**”

Neural Network origins

- 1958 **Frank Rosenblatt** (psychologist) proposes the classic **perceptron** (improving on 1943 McCulloch & Pitts neural model)

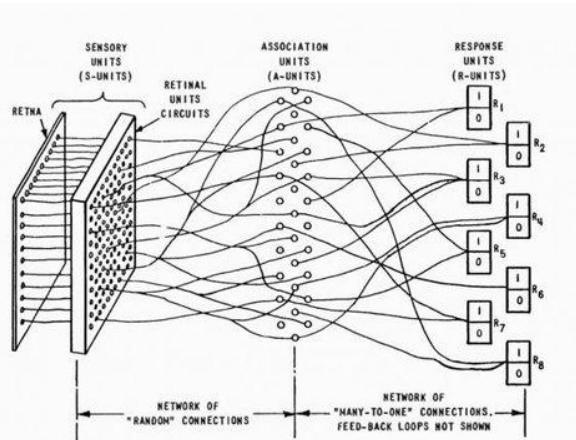
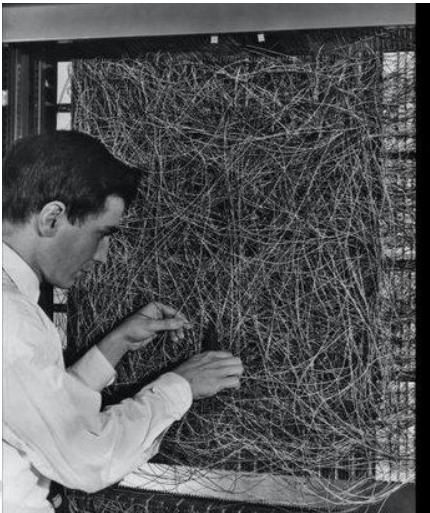


Figure 1 ORGANIZATION OF THE MARK I PERCEPTRON

Neural Network origins

- 1954 **Software perceptron: IBM 704** 1st mass produced floating point computer (Fortran, LISP...)




Perceptron: a model neuron

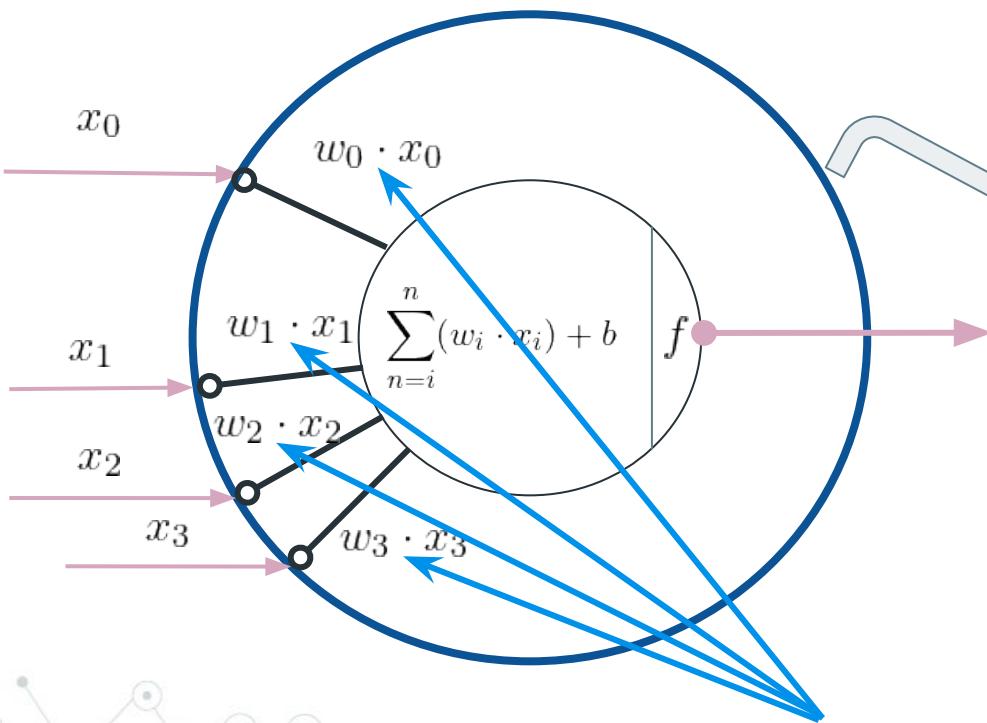
Santiago Ramón y Cajal
(1889): Neurons are cells =
individual units
communicate by synapsis



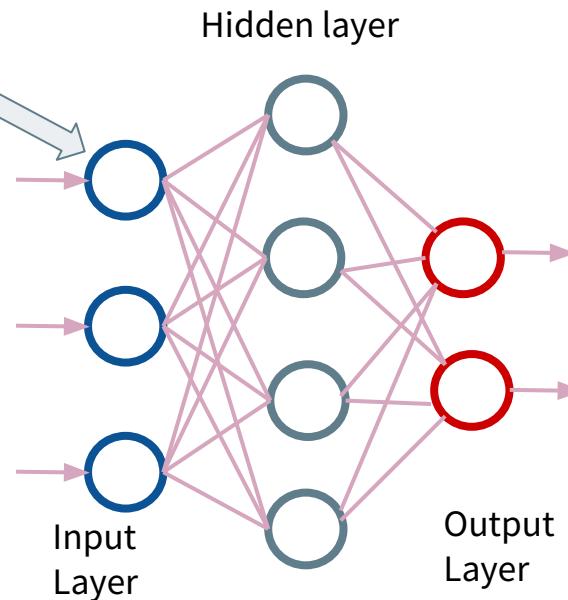
Perceptron: a model neuron



Artificial Neural Network: Multi-layer perceptron

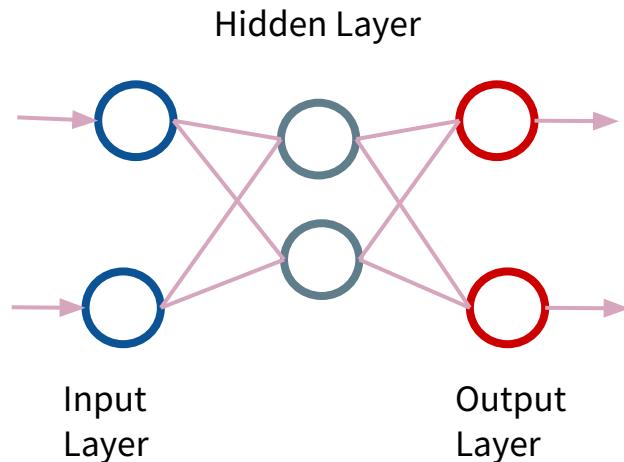


Learn = Neurons adjust their weights

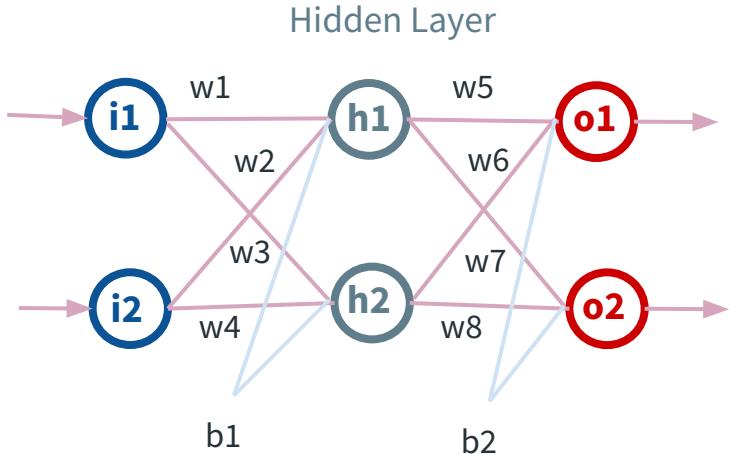


Activation function

- Without a **non-linear activation function (f)** the neural network can only account for linear effects



Activation function (linear)




$$h1 = w1 \cdot i1 + w2 \cdot i2 + b1$$

$$h2 = w3 \cdot i1 + w4 \cdot i2 + b1$$

$$o1 = w5 \cdot h1 + w6 \cdot h2 + b2$$

$$o2 = w7 \cdot h1 + w8 \cdot h2 + b2$$

Activation function (linear)



$$o_1 = w_5 \cdot (w_1 \cdot i_1 + w_2 \cdot i_2 + b_1) + w_6 \cdot (w_3 \cdot i_1 + w_4 \cdot i_2 + b_1) + b_2$$

$$o_1 = w_5 \cdot w_1 \cdot i_1 + w_5 \cdot w_2 \cdot i_2 + w_5 \cdot b_1 + w_6 \cdot w_3 \cdot i_1 + w_6 \cdot w_4 \cdot i_2 + w_6 \cdot b_1 + b_2$$

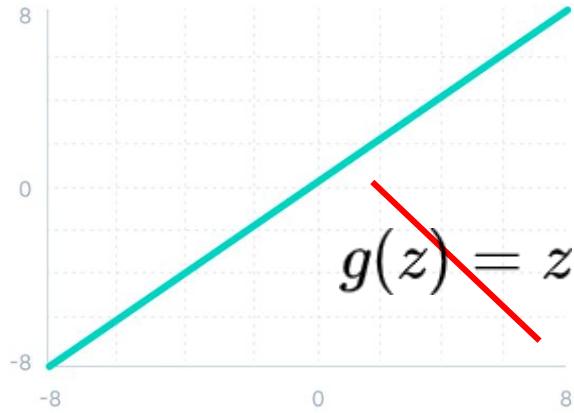
$$o_1 = (w_5 \cdot w_1 + w_6 \cdot w_3) \cdot i_1 + (w_5 \cdot w_2 + w_6 \cdot w_4) \cdot i_2 + (w_5 \cdot b_1 + w_6 \cdot b_1 + b_2)$$

$$o_1 = A_1 \cdot i_1 + A_2 \cdot i_2 + C_1$$


Linear again!

Activation Functions

Linear (no!)



Sigmoid / Softmax*

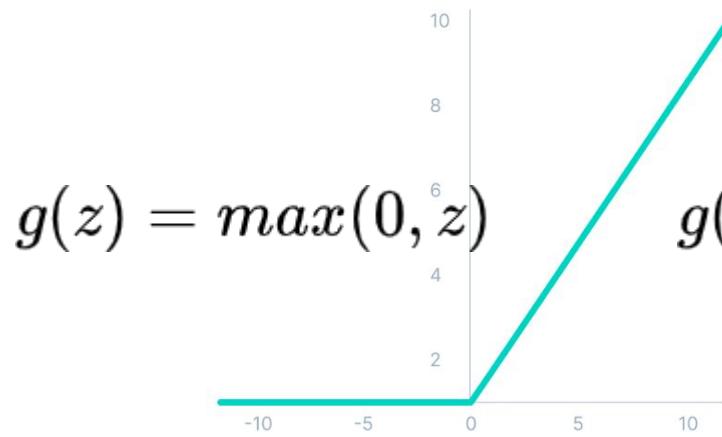


$$g(z) = \frac{1}{1+e^{-z}}$$

$$* g(z_i) = \frac{e^{z_i}}{\sum e^{z_j}}$$

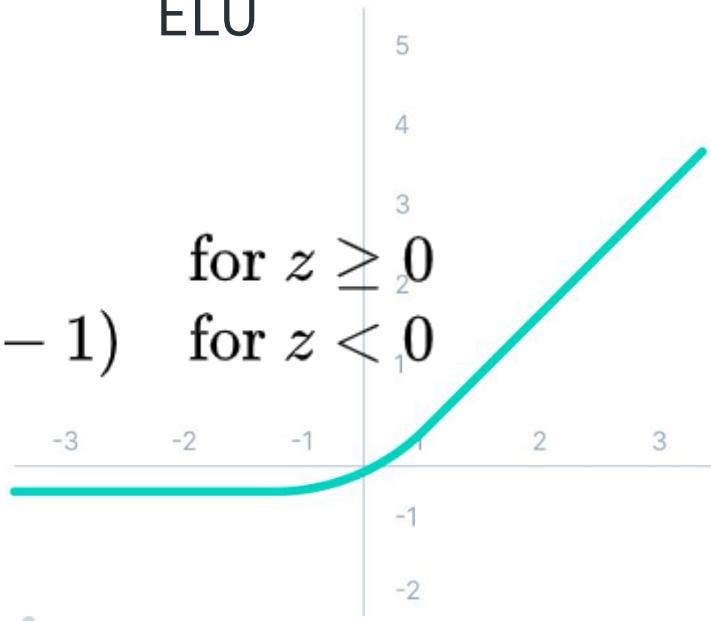
Activation Functions

ReLU



ELU

$$g(z) = \begin{cases} z & \text{for } z \geq 0 \\ \alpha(e^z - 1) & \text{for } z < 0 \end{cases}$$



Universal Approximation Theorem

Why we **can** use a NN:

“For any **continuous function** for a hypercube $[0,1]^d$ to real numbers, and every positive epsilon, there exists a **sigmoid** based **1-HIDDEN LAYER NEURAL NETWORK** that obtains at most epsilon error in functional space” Cybenko '89

→ Big enough **NN** can *approximate* (not represent) any smooth function



Universal Approximation Theorem

Why we **can** use a NN:


“For any **continuous function** for a hypercube $[0,1]^d$ to real numbers, **non-constant, bounded and continuous activation function f**, and every positive epsilon, there exists **1-HIDDEN LAYER NEURAL NETWORK** using f that obtains at most epsilon error in functional space”

Horvik '91

→ Big enough **NN** can *approximate* (not represent) any smooth function




Universal Approximation Theorem (Intuition)




deepmind
@Czarnecki

Universal Approximation Theorem (Intuition)



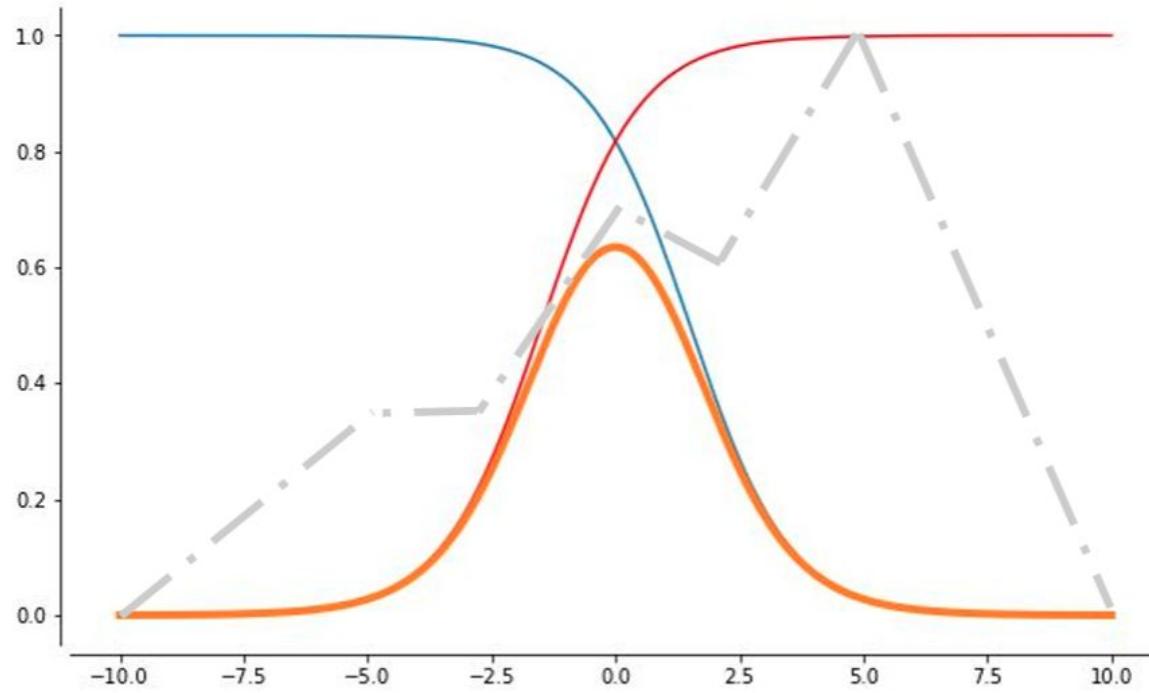
deepmind
@Czarnecki

Universal Approximation Theorem (Intuition)




deepmind
@Czarnecki

Universal Approximation Theorem (Intuition)




deepmind
@Czarnecki

Universal Approximation Theorem (Intuition)




deepmind
@Czarnecki

Universal Approximation Theorem (Intuition)



deepmind
@Czarnecki


Artificial Neural Network: Multi-layer perceptron




$$f(\vec{x}) = g(W \cdot \vec{x} + \vec{b})$$

W is an array,
b is a vector

DEEP Learning: many hidden layers




DEEP Learning: many hidden layers




Hidden layer

$$h(x) = g(z^h(x)) = g(W^h \cdot x + b^h)$$

DEEP Learning: many hidden layers



DEEP Learning: many hidden layers



Minimize the loss

- ◎ Find the weights that generate **minimum loss** (an arbitrary multi-dimensional function!)

Simplify: Taylor approximation and use standard algorithms...

Minimize the loss

(1st derivative) **Gradient Descent**

$$W_{t+1} = W_t - \lambda_t \nabla f(W_t)$$

Newton (2nd derivative)

$$W_{t+1} = W_t - [Hf(W_t)]^{-1} \nabla f(W_t)$$

hessian

$$H = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$



Minimize the loss

(1st derivative) **Gradient Descent**

$$W_{t+1} = W_t - \lambda_t \nabla f(W_t)$$

learning rate

gradient



Newton is *fast* BUT expensive

- And not always works (smooth functions)

Newton (2nd derivative)



$$W_{t+1} = W_t - [Hf(W_t)]^{-1} \nabla f(W_t)$$

hessian


$$H = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$



Gradient Descent




Backpropagation (neuron level)




Credits: A. Karpathy

Backpropagation (neuron level)




Credits: A. Karpathy

Backpropagation (neuron level)




Credits: A. Karpathy

Backpropagation (neuron level)



Backpropagation (neuron level)



Backpropagation (neuron level)

x

Activations

$$\frac{\partial L}{\partial x}$$

$$= \frac{\partial L}{\partial z} \frac{\partial z}{\partial x}$$

y

$$\frac{\partial L}{\partial y} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial y}$$

Local gradient

$$\frac{\partial z}{\partial x}$$

$$\frac{\partial z}{\partial y}$$

f

z

$$\frac{\partial L}{\partial z}$$


Gradients
(backpropagation)

Credits: A. Karpathy


Fit linear NN / non-linear NN

https://github.com/cwestend/IACDEEP_introNN







Fit linear NN / non-linear NN



Fit linear NN / non-linear NN




Fit linear NN / non-linear NN







Fit linear NN / non-linear NN



Three Models: 1D Poly, Linear NN, Non-linear NN





Fit linear NN / non-linear NN with splits (80/10/10)



Non-Linear Network Architecture



Fit linear NN / non-linear NN with splits (80/10/10)



Takeaways:

- **Deep learning**: uses ANN (many hidden), learns from data
- Needs **non-linear** activation functions
- Minimize **loss** (learn):
 - *gradient descent, learning rate, backpropagation*
- **Normalize** data (**must**)
- Needs **lots of data!**

https://github.com/cwestend/IACDEEP_introNN

